

Model Persistence - Loading

Auxiliary materials

```
In [1]: import numpy as np
import torch
import lightning.pytorch as pl
import torchmetrics
import torchvision
from torchinfo import summary
from torchview import draw_graph
import pandas as pd
import matplotlib.pyplot as plt
```

Iris Net

```
In [2]: class MultiClassModule(pl.LightningModule):
    def __init__(self,
                 output_size,
                 **kwargs):
        super().__init__(**kwargs)
        self.mc_acc = torchmetrics.classification.Accuracy(task='multiclass',
                                                            num_classes=output_size)
        self.cce_loss = torch.nn.CrossEntropyLoss()

    def predict(self, x):
        return torch.softmax(self(x), -1)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=0.001)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        x, y_true = train_batch
        y_pred = self(x)
        perm = (0, -1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_true)
        loss = self.cce_loss(y_pred.permute(*perm), y_true)
        self.log('train_acc', acc, on_step=False, on_epoch=True)
        self.log('train_loss', loss, on_step=False, on_epoch=True)
        return loss

    def validation_step(self, val_batch, batch_idx):
        x, y_true = val_batch
        y_pred = self(x)
        perm = (0, -1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_true)
        loss = self.cce_loss(y_pred.permute(*perm), y_true)
        self.log('val_acc', acc, on_step=False, on_epoch=True)
        self.log('val_loss', loss, on_step=False, on_epoch=True)
        return loss

    def test_step(self, test_batch, batch_idx):
        x, y_true = test_batch
```

```

y_pred = self(x)
perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
acc = self.mc_acc(y_pred.permute(*perm),y_true)
loss = self.cce_loss(y_pred.permute(*perm),y_true)
self.log('test_acc', acc, on_step=False, on_epoch=True)
self.log('test_loss', loss, on_step=False, on_epoch=True)
return loss

```

```

In [3]: class MultiLayerNetwork(MultiClassModule):
        def __init__(self,
                      input_size,
                      output_size,
                      hidden_size = 64,
                      **kwargs):
            super().__init__(output_size=output_size,
                             **kwargs)
            self.save_hyperparameters()
            self.hidden = torch.nn.Linear(input_size,
                                           hidden_size)
            self.output = torch.nn.Linear(hidden_size,
                                           output_size)

        def forward(self, x):
            y = x
            y = torch.tanh(self.hidden(y))
            y = self.output(y)
            return y

```

```

In [4]: data = np.loadtxt("https://www.cs.mtsu.edu/~jphillips/courses/CSCI4850-5850/publ
np.random.seed(0)
np.random.shuffle(data)
X = data[:, :-1]
Y = data[:, -1]
split_point = int(X.shape[0] * 0.8)
# The dataLoaders handle shuffling, batching, etc...
xy_train = torch.utils.data.DataLoader(list(zip(torch.tensor(X[:split_point]).floa
                                                torch.tensor(Y[:split_point]).long
                                                shuffle=True, batch_size=32,num_workers=4
xy_val = torch.utils.data.DataLoader(list(zip(torch.tensor(X[split_point:]).floa
                                                torch.tensor(Y[split_point:]).long
                                                shuffle=False, batch_size=32,num_workers=

```

/opt/conda/lib/python3.11/site-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested maximum number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn(_create_warning_msg(
```

```

In [5]: iris_net = MultiLayerNetwork.load_from_checkpoint("iris_net.ckpt")

# Or use the one made by Lightning automatically...
# iris_net = MultiLayerNetwork.load_from_checkpoint("logs/persistence/iris_net-0

```

```

In [6]: x = torch.Tensor(X[0:1]).to(iris_net.device)
x

```

```

Out[6]: tensor([[6.4000, 2.9000, 4.3000, 1.3000]])

```

```
In [7]: iris_net.predict(x).cpu().detach().numpy()
```

```
Out[7]: array([[0.02313578, 0.9282432 , 0.048621  ]], dtype=float32)
```

```
In [8]: Y[0:1]
```

```
Out[8]: array([1.])
```

Just validation...

```
In [9]: logger = pl.loggers.CSVLogger("logs",
                                     name="persistence")
```

```
In [10]: trainer = pl.Trainer(logger=logger,
                              max_epochs=100,
                              enable_progress_bar=True,
                              log_every_n_steps=0,
                              enable_checkpointing=True, # Notice this!
                              callbacks=[pl.callbacks.TQDMProgressBar(refresh_rate=50)])
```

```
GPU available: True (cuda), used: True
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs
```

This result should match the previous result which indicates successful loading of the saved model.

```
In [11]: trainer.validate(iris_net, xy_val)
```

```
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-requeueing enabled. Setting signal handlers.
Validation: |          | 0/? [00:00<?, ?it/s]
```

Validate metric	DataLoader 0
val_acc	1.0
val_loss	0.12450059503316879

```
Out[11]: [{'val_acc': 1.0, 'val_loss': 0.12450059503316879}]
```

Encoder-Decoder

```
In [12]: # Tokenization Tools
import os
os.environ["TOKENIZERS_PARALLELISM"] = "false"
import tokenizers
import io
```

```
In [13]: import urllib
data = []
my_url = "https://raw.githubusercontent.com/luisroque/deep-learning-articles/mai
with urllib.request.urlopen(my_url) as raw_data:
```

```
for line in raw_data:
    data.append(line.decode("utf-8").split('\t')[0:2])
data = np.array(data)
```

In [14]: data.shape

Out[14]: (170304, 2)

```
In [15]: # Subset? - All of the data will take some time...
n_seq = data.shape[0]
n_seq = 10000
data = data[0:n_seq]
split_point = int(data.shape[0] * 0.8) # Keep 80/20 split
np.random.seed(0)
np.random.shuffle(data) # In-place modification
```

In [16]: data[0]

Out[16]: array(['These are real.', 'Estas são autênticas.'], dtype='<U184')

```
In [17]: eng = data[:,0]
por = data[:,1]
```

In [18]: eng.shape

Out[18]: (10000,)

In [19]: eng[0:5]

Out[19]: array(['These are real.', "I'm sorry.", 'I have wine.', "I won't do it.",
 'I eat bread.'], dtype='<U184')

In [20]: por.shape

Out[20]: (10000,)

In [21]: por[0:5]

Out[21]: array(['Estas são autênticas.', 'Desculpe!', 'Tenho vinho.',
 'Eu não irei fazer isso.', 'Eu como pão.'], dtype='<U184')

Pretrained tokenizer from the persisted JSON file

```
In [22]: eng_tokenizer = tokenizers.Tokenizer(tokenizers.models.BPE()).from_file("eng_tra
por_tokenizer = tokenizers.Tokenizer(tokenizers.models.BPE()).from_file("por_tra
```

In [23]: eng_tokenizer.encode("Here is a test. ").tokens

Out[23]: ['Here', 'is', 'a', 'test', '.']

```
In [24]: temp = eng_tokenizer.encode("Here is a test. ").ids
temp
```

Out[24]: [443, 78, 46, 3165, 9]

```
In [25]: temp = eng_tokenizer.decode(temp + [0,0])
temp
```

```
Out[25]: 'Here is a test .'
```

```
In [26]: eng_tokenizer.get_vocab_size()
```

```
Out[26]: 3694
```

```
In [27]: por_tokenizer.get_vocab_size()
```

```
Out[27]: 6459
```

```
In [28]: eng_recoded = np.array([eng_tokenizer.decode(eng_tokenizer.encode(s).ids) for s
por_recoded = np.array([por_tokenizer.decode(por_tokenizer.encode(s).ids) for s
```

```
In [29]: eng_recoded[0]
```

```
Out[29]: 'These are real .'
```

```
In [30]: eng[0]
```

```
Out[30]: 'These are real.'
```

```
In [31]: por_recoded[0]
```

```
Out[31]: 'Estas são autênticas .'
```

```
In [32]: por[0]
```

```
Out[32]: 'Estas são autênticas.'
```

```
In [33]: def encode_seq(x,tokenizer,max_length=0):
# String to integer
x = tokenizer.encode("<START>"+x+"<STOP>").ids
x += [0]*(max_length-len(x))
return x

def decode_seq(x,tokenizer):
return tokenizer.decode(x)
```

```
In [34]: temp = encode_seq(eng_recoded[0],eng_tokenizer,20)
temp
```

```
Out[34]: [1, 425, 140, 442, 9, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
In [35]: len(temp)
```

```
Out[35]: 20
```

```
In [36]: decode_seq(temp,eng_tokenizer)
```

```
Out[36]: 'These are real .'
```

```
In [37]: temp = encode_seq(por_recoded[0],por_tokenizer,20)
```

```
temp
```

```
Out[37]: [1, 862, 229, 6063, 8, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
In [38]: decode_seq(temp,por_tokenizer)
```

```
Out[38]: 'Estas são autênticas .'
```

```
In [39]: max_eng = np.max([len(encode_seq(i,eng_tokenizer)) for i in eng_recoded])
max_eng
```

```
Out[39]: 12
```

```
In [40]: max_por = np.max([len(encode_seq(i,por_tokenizer)) for i in por_recoded])
max_por
```

```
Out[40]: 11
```

```
In [41]: X = np.vstack([encode_seq(x,eng_tokenizer,max_eng) for x in eng_recoded])
Y = np.vstack([encode_seq(x,por_tokenizer,max_por) for x in por_recoded])
```

```
In [42]: enc_x_train = X[:split_point]
enc_x_val = X[split_point:]
enc_x_train
```

```
Out[42]: array([[ 1, 425, 140, ..., 0, 0, 0],
 [ 1, 30, 6, ..., 0, 0, 0],
 [ 1, 30, 152, ..., 0, 0, 0],
 ...,
 [ 1, 90, 6, ..., 0, 0, 0],
 [ 1, 30, 6, ..., 0, 0, 0],
 [ 1, 1592, 21, ..., 0, 0, 0]])
```

```
In [43]: dec_x_train = Y[:,0:-1][:split_point]
dec_x_val = Y[:,0:-1][split_point:]
dec_x_train
```

```
Out[43]: array([[ 1, 862, 229, ..., 0, 0, 0],
 [ 1, 3279, 3, ..., 0, 0, 0],
 [ 1, 352, 719, ..., 0, 0, 0],
 ...,
 [ 1, 141, 416, ..., 0, 0, 0],
 [ 1, 188, 850, ..., 0, 0, 0],
 [ 1, 3776, 19, ..., 0, 0, 0]])
```

```
In [44]: dec_y_train = Y[:,1:][:split_point]
dec_y_val = Y[:,1:][split_point:]
dec_y_train
```

```
Out[44]: array([[ 862, 229, 6063, ..., 0, 0, 0],
 [3279, 3, 2, ..., 0, 0, 0],
 [ 352, 719, 8, ..., 0, 0, 0],
 ...,
 [ 141, 416, 3597, ..., 0, 0, 0],
 [ 188, 850, 8, ..., 0, 0, 0],
 [3776, 19, 2, ..., 0, 0, 0]])
```

```
In [45]: print(enc_x_train.shape)
print(dec_x_train.shape)
```

```
print(dec_y_train.shape)
```

```
(8000, 12)
(8000, 10)
(8000, 10)
```

```
In [46]: print(enc_x_val.shape)
         print(dec_x_val.shape)
         print(dec_y_val.shape)
```

```
(2000, 12)
(2000, 10)
(2000, 10)
```

```
In [47]: batch_size = 256
         xy_train = torch.utils.data.DataLoader(list(zip(torch.Tensor(enc_x_train).long(),
                                                         torch.Tensor(dec_x_train).long(),
                                                         torch.Tensor(dec_y_train).long(),
                                                         shuffle=True, batch_size=batch_size,
                                                         num_workers=4)
         xy_val = torch.utils.data.DataLoader(list(zip(torch.Tensor(enc_x_val).long(),
                                                         torch.Tensor(dec_x_val).long(),
                                                         torch.Tensor(dec_y_val).long()),
                                                         shuffle=False, batch_size=batch_size,
                                                         num_workers=4)
```

Model Loading...

```
In [48]: class RecurrentResidual(pl.LightningModule):
         def __init__(self,
                     latent_size = 256,
                     bidirectional = False,
                     **kwargs):
             super().__init__(**kwargs)
             self.layer_norm = torch.nn.LayerNorm(latent_size)
             self.rnn_layer = torch.nn.LSTM(latent_size,
                                             latent_size // 2 if bidirectional else la
                                             bidirectional=bidirectional,
                                             batch_first=True)

         def forward(self, x):
             return x + self.rnn_layer(self.layer_norm(x))[0]
```

```
In [49]: class EncoderNetwork(pl.LightningModule):
         def __init__(self,
                     num_tokens,
                     latent_size = 256, # Use something divisible by 2
                     n_layers = 8,
                     **kwargs):
             super().__init__(**kwargs)
             self.embedding = torch.nn.Embedding(num_tokens,
                                                  latent_size,
                                                  padding_idx=0)

             self.dropout = torch.nn.Dropout1d(0.05) # Whole token dropped
             self.rnn_layers = torch.nn.Sequential(*[
                 RecurrentResidual(latent_size, True) for _ in range(n_layers)
             ])

         def forward(self, x):
             y = x
```

```

y = self.embedding(y)
y = self.dropout(y)
y = self.rnn_layers(y)[:,-1]
return y

```

```

In [50]: class DecoderNetwork(pl.LightningModule):
def __init__(self,
            num_tokens,
            latent_size = 256, # Use something divisible by 2
            n_layers = 8,
            **kwargs):
    super().__init__(**kwargs)
    self.embedding = torch.nn.Embedding(num_tokens,
                                       latent_size,
                                       padding_idx=0)
    # self.dropout = torch.nn.Dropout1d(0.1) # Whole token dropped
    self.linear = torch.nn.Linear(latent_size*2,
                                  latent_size)
    self.rnn_layers = torch.nn.Sequential(*[
        RecurrentResidual(latent_size,False) for _ in range(n_layers)
    ])
    self.output_layer = torch.nn.Linear(latent_size,
                                       num_tokens)

def forward(self, x_enc, x_dec):
    y_enc = x_enc.unsqueeze(1).repeat(1,x_dec.shape[1],1)
    y_dec = self.embedding(x_dec)
    # y_dec = self.dropout(y_dec)
    y = y_enc
    y = torch.concatenate([y_enc,y_dec],-1)
    y = self.linear(y)
    y = self.rnn_layers(y)
    y = self.output_layer(y)
    return y

```

```

In [51]: class EncDecLightningModule(pl.LightningModule):
def __init__(self,
            output_size,
            **kwargs):
    super().__init__(**kwargs)
    self.mc_acc = torchmetrics.classification.Accuracy(task='multiclass',
                                                       num_classes=output_si
                                                       ignore_index=0)
    self.cce_loss = torch.nn.CrossEntropyLoss(ignore_index=0)

def predict(self, x):
    return torch.softmax(self(x),-1)

def configure_optimizers(self):
    optimizer = torch.optim.Adam(self.parameters(), lr=0.001)
    return optimizer

def training_step(self, train_batch, batch_idx):
    x_enc, x_dec, y_dec = train_batch
    y_pred = self(x_enc, x_dec)
    perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
    acc = self.mc_acc(y_pred.permute(*perm),y_dec)
    loss = self.cce_loss(y_pred.permute(*perm),y_dec)
    self.log('train_acc', acc, on_step=False, on_epoch=True)

```



```

        self.log('train_loss', loss, on_step=False, on_epoch=True)
        return loss

    # Validate used for Teacher Forcing
    def validation_step(self, val_batch, batch_idx):
        x_enc, x_dec, y_dec = val_batch
        y_pred = self(x_enc, x_dec)
        perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_dec)
        loss = self.cce_loss(y_pred.permute(*perm), y_dec)
        self.log('val_acc', acc, on_step=False, on_epoch=True)
        self.log('val_loss', loss, on_step=False, on_epoch=True)
        return loss

    # Test used for Non-Teacher Forcing
    def test_step(self, test_batch, batch_idx):
        x_enc, x_dec, y_dec = test_batch
        context = self.enc_net(x_enc)
        tokens = torch.zeros_like(x_dec).long()
        tokens[:,0] = 1
        for i in range(y_dec.shape[1]-1):
            tokens[:,i+1] = self.dec_net(context, tokens).argmax(-1)[: ,i]
        y_pred = self(x_enc, tokens)
        perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_dec)
        loss = self.cce_loss(y_pred.permute(*perm), y_dec)
        self.log('test_acc', acc, on_step=False, on_epoch=True)
        self.log('test_loss', loss, on_step=False, on_epoch=True)
        return loss

    def predict_step(self, predict_batch, batch_idx):
        x_enc, x_dec, y_dec = test_batch
        context = self.enc_net(x_enc)
        tokens = torch.zeros_like(x_dec).long()
        tokens[:,0] = 1
        for i in range(y_dec.shape[1]-1):
            tokens[:,i+1] = self.dec_net(context, tokens).argmax(-1)[: ,i]
        y_pred = self(x_enc, tokens)
        return y_pred

```

```

In [52]: class EncDecNetwork(EncDecLightningModule):
    def __init__(self,
                  num_enc_tokens,
                  num_dec_tokens,
                  latent_size = 256, # Use something divisible by 2
                  n_layers = 8,
                  **kwargs):
        super().__init__(output_size=num_dec_tokens,
                          **kwargs)
        self.save_hyperparameters()
        self.enc_net = EncoderNetwork(num_enc_tokens, latent_size, n_layers)
        self.dec_net = DecoderNetwork(num_dec_tokens, latent_size, n_layers)

    def forward(self, x_enc, x_dec):
        return self.dec_net(self.enc_net(x_enc), x_dec)

```

```

In [53]: enc_dec_net = EncDecNetwork.load_from_checkpoint("enc_dec_net.ckpt")

```

Test without Teacher Forcing

```
In [54]: # Complete max_length cycles with the decoder
i = 0
enc_dec_net.to("cpu")
context = enc_dec_net.enc_net(torch.Tensor(enc_x_val[i:i+1]).long())
token = torch.zeros((1,dec_y_val.shape[1])).long()
token[0,0] = 1

for x in range(dec_y_val.shape[1]-1):
    result = enc_dec_net.dec_net(context,token).argmax(-1)
    token[0,x+1] = result[0,x]
    if result[0,x] == 2:
        break
result = token.cpu().detach().numpy()[0]
result
```

```
Out[54]: array([ 1, 5549, 514, 19, 2, 0, 0, 0, 0])
```

English input...

```
In [55]: decode_seq(enc_x_val[i],eng_tokenizer)
```

```
Out[55]: 'Can you swim ?'
```

Portuguese translation from network...

```
In [56]: decode_seq(result,por_tokenizer)
```

```
Out[56]: 'Sabes nadar ?'
```

Target translation from the data set...

```
In [57]: decode_seq(dec_y_val[i],por_tokenizer)
```

```
Out[57]: 'Sabe nadar ?'
```

```
In [58]: result.shape
```

```
Out[58]: (10,)
```

```
In [59]: dec_y_val.shape
```

```
Out[59]: (2000, 10)
```

Accuracy **without** teacher forcing...

```
In [60]: logger = pl.loggers.CSVLogger("logs",
                                         name="persistence")
```

```
In [61]: trainer = pl.Trainer(logger=logger,
                               max_epochs=30,
                               enable_progress_bar=True,
                               log_every_n_steps=0,
```

```
enable_checkpointing=True, # Notice this here!
callbacks=[pl.callbacks.TQDMProgressBar(refresh_rate=50)])
```

```
GPU available: True (cuda), used: True
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs
```

This result should match the previous result which indicates successful loading of the saved model.

```
In [62]: trainer.test(enc_dec_net, xy_val)
```

```
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-requeueing enabled. Setting signal handlers.
Testing: |          | 0/? [00:00<?, ?it/s]
```

Test metric	DataLoader 0
test_acc	0.4213317930698395
test_loss	6.924007892608643

```
Out[62]: [{'test_acc': 0.4213317930698395, 'test_loss': 6.924007892608643}]
```

```
In [ ]:
```