

Model Persistence - Saving

Auxiliary materials

```
In [1]: import numpy as np
import torch
import lightning.pytorch as pl
import torchmetrics
import torchvision
from torchinfo import summary
from torchview import draw_graph
import pandas as pd
import matplotlib.pyplot as plt
```

Iris Net

```
In [2]: class MultiClassModule(pl.LightningModule):
    def __init__(self,
                 output_size,
                 **kwargs):
        super().__init__(**kwargs)
        self.mc_acc = torchmetrics.classification.Accuracy(task='multiclass',
                                                            num_classes=output_size)
        self.cce_loss = torch.nn.CrossEntropyLoss()

    def predict(self, x):
        return torch.softmax(self(x), -1)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=0.001)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        x, y_true = train_batch
        y_pred = self(x)
        perm = (0, -1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_true)
        loss = self.cce_loss(y_pred.permute(*perm), y_true)
        self.log('train_acc', acc, on_step=False, on_epoch=True)
        self.log('train_loss', loss, on_step=False, on_epoch=True)
        return loss

    def validation_step(self, val_batch, batch_idx):
        x, y_true = val_batch
        y_pred = self(x)
        perm = (0, -1) + tuple(range(y_pred.ndim))[1:-1]
        acc = self.mc_acc(y_pred.permute(*perm), y_true)
        loss = self.cce_loss(y_pred.permute(*perm), y_true)
        self.log('val_acc', acc, on_step=False, on_epoch=True)
        self.log('val_loss', loss, on_step=False, on_epoch=True)
        return loss

    def test_step(self, test_batch, batch_idx):
        x, y_true = test_batch
```

```

y_pred = self(x)
perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
acc = self.mc_acc(y_pred.permute(*perm),y_true)
loss = self.cce_loss(y_pred.permute(*perm),y_true)
self.log('test_acc', acc, on_step=False, on_epoch=True)
self.log('test_loss', loss, on_step=False, on_epoch=True)
return loss

```

```

In [3]: data = np.loadtxt("https://www.cs.mtsu.edu/~jphillips/courses/CSCI4850-5850/publ
np.random.seed(0)
np.random.shuffle(data)
X = data[:, :-1]
Y = data[:, -1]
split_point = int(X.shape[0] * 0.8)
# The dataLoaders handle shuffling, batching, etc...
xy_train = torch.utils.data.DataLoader(list(zip(torch.tensor(X[:split_point]).floa
torch.tensor(Y[:split_point]).long
shuffle=True, batch_size=32,num_workers=4
xy_val = torch.utils.data.DataLoader(list(zip(torch.tensor(X[split_point:]).floa
torch.tensor(Y[split_point:]).long
shuffle=False, batch_size=32,num_workers=

```

/opt/conda/lib/python3.11/site-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested maximum number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

warnings.warn(_create_warning_msg(

```

In [4]: class MultiLayerNetwork(MultiClassModule):
def __init__(self,
input_size,
output_size,
hidden_size = 64,
**kwargs):
super().__init__(output_size=output_size,
**kwargs)
self.save_hyperparameters()
self.hidden = torch.nn.Linear(input_size,
hidden_size)
self.output = torch.nn.Linear(hidden_size,
output_size)

def forward(self, x):
y = x
y = torch.tanh(self.hidden(y))
y = self.output(y)
return y

```

```

In [5]: iris_net = MultiLayerNetwork(X.shape[1],
len(np.unique(Y)))

```

```

In [6]: x = torch.Tensor(X[0:1]).to(iris_net.device)
x

```

```

Out[6]: tensor([[6.4000, 2.9000, 4.3000, 1.3000]])

```

```

In [7]: iris_net.predict(x).cpu().detach().numpy()

```

Out[7]: array([[0.22635038, 0.43114087, 0.34250873]], dtype=float32)

In [8]: Y[1]

Out[8]: 0.0

In [9]: logger = pl.loggers.CSVLogger("logs",
name="persistence",
version="iris_net-0")

In [10]: trainer = pl.Trainer(logger=logger,
max_epochs=100,
enable_progress_bar=True,
log_every_n_steps=0,
enable_checkpointing=True, # Notice this!
callbacks=[pl.callbacks.TQDMProgressBar(refresh_rate=50)])

GPU available: True (cuda), used: True
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs

In [12]: trainer.validate(iris_net, xy_val)

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-requeueing enabled. Setting signal handlers.
Validation: | | 0/? [00:00<?, ?it/s]

Validate metric	DataLoader 0
val_acc	0.2666666805744171
val_loss	1.0410279035568237

Out[12]: [{'val_acc': 0.2666666805744171, 'val_loss': 1.0410279035568237}]

In [13]: trainer.fit(iris_net, xy_train, xy_val)

/opt/conda/lib/python3.11/site-packages/lightning/pytorch/callbacks/model_checkpoint.py:639: Checkpoint directory logs/persistence/iris_net-0/checkpoints exists and is not empty.
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

	Name	Type	Params
0	mc_acc	MulticlassAccuracy	0
1	cce_loss	CrossEntropyLoss	0
2	hidden	Linear	320
3	output	Linear	195

515 Trainable params
0 Non-trainable params
515 Total params
0.002 Total estimated model params size (MB)
SLURM auto-requeueing enabled. Setting signal handlers.
Sanity Checking: | | 0/? [00:00<?, ?it/s]
Training: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]
Validation: | | 0/? [00:00<?, ?it/s]

[illegible]

```
`Trainer.fit` stopped: `max epochs=100` reached.
```

```
In [14]: trainer.validate(iris_net, xy_val)
```

Validate metric	DataLoader 0
val_acc	1.0
val_loss	0.12450060248374939

```
In [18]: x = torch.Tensor(X[0:1]).to(iris_net.device)
          x
```

```
Out[18]: tensor([[6.4000, 2.9000, 4.3000, 1.3000]])
```

```
In [19]: iris_net.predict(x).cpu().detach().numpy()
```

```
Out[19]: array([[0.02313578, 0.9282432 , 0.048621  ]], dtype=float32)
```

```
In [20]: Y[0:1]
```

```
Out[20]: array([1.])
```

Here's the magic word:

```
In [21]: trainer.save_checkpoint("iris_net.ckpt")
```

Switch to Loading example notebook...

Encoder-Decoder

```
In [22]: # Tokenization Tools
import os
os.environ["TOKENIZERS_PARALLELISM"] = "false"
import tokenizers
import io
```

```
In [23]: import urllib
data = []
my_url = "https://raw.githubusercontent.com/luisroque/deep-learning-articles/main/
with urllib.request.urlopen(my_url) as raw_data:
    for line in raw_data:
        data.append(line.decode("utf-8").split('\t')[0:2])
data = np.array(data)
```

```
In [24]: data.shape
```

```
Out[24]: (170304, 2)
```

```
In [25]: # Subset? - All of the data will take some time...
n_seq = data.shape[0]
n_seq = 10000
data = data[0:n_seq]
split_point = int(data.shape[0] * 0.8) # Keep 80/20 split
np.random.seed(0)
np.random.shuffle(data) # In-place modification
```

```
In [26]: data[0]
```

```
Out[26]: array(['These are real.', 'Estas são autênticas.'], dtype='<U184')
```

```
In [27]: eng = data[:,0]
por = data[:,1]
```

```
In [28]: eng.shape
```

Out[28]: (10000,)

In [29]: eng[0:5]

Out[29]: array(['These are real.', 'I'm sorry.', 'I have wine.', 'I won't do it.',
'I eat bread.'], dtype='<U184')

In [30]: por.shape

Out[30]: (10000,)

In [31]: por[0:5]

Out[31]: array(['Estas são autênticas.', 'Desculpe!', 'Tenho vinho.',
'Eu não irei fazer isso.', 'Eu como pão.'], dtype='<U184')

```
In [32]: unknown_token = "<UNK>" # token for unknown words
special_tokens = [unknown_token, "<START>", "<STOP>"] # special tokens

eng_tokenizer = tokenizers.Tokenizer(tokenizers.models.BPE(unk_token=unknown_tok
eng_token_trainer = tokenizers.trainers.BpeTrainer(vocab_size=100000,special_tok
eng_tokenizer.pre_tokenizer = tokenizers.pre_tokenizers.Whitespace()

por_tokenizer = tokenizers.Tokenizer(tokenizers.models.BPE(unk_token=unknown_tok
por_token_trainer = tokenizers.trainers.BpeTrainer(vocab_size=100000,special_tok
por_tokenizer.pre_tokenizer = tokenizers.pre_tokenizers.Whitespace()
```

```
In [33]: with open("eng_strings.txt","w") as f:
        for s in eng:
            f.write(s)
            f.write("\n")
        with open("por_strings.txt","w") as f:
            for s in por:
                f.write(s)
                f.write("\n")
```

```
In [34]: eng_tokenizer.train(["eng_strings.txt"],eng_token_trainer)
por_tokenizer.train(["por_strings.txt"],por_token_trainer)
```

Persist the tokenizer...

```
In [35]: eng_tokenizer.save("eng_trained.json")
por_tokenizer.save("por_trained.json")
```

```
In [36]: eng_tokenizer.encode("Here is a test.").tokens
```

Out[36]: ['Here', 'is', 'a', 'test', '.']

```
In [37]: temp = eng_tokenizer.encode("Here is a test.").ids
temp
```

Out[37]: [443, 78, 46, 3165, 9]

```
In [38]: temp = eng_tokenizer.decode(temp + [0,0])
temp
```

Out[38]: 'Here is a test .'

```
In [39]: eng_tokenizer.get_vocab_size()
```

Out[39]: 3694

```
In [40]: por_tokenizer.get_vocab_size()
```

Out[40]: 6459

```
In [41]: eng_recoded = np.array([eng_tokenizer.decode(eng_tokenizer.encode(s).ids) for s
por_recoded = np.array([por_tokenizer.decode(por_tokenizer.encode(s).ids) for s
```

```
In [42]: eng_recoded[0]
```

Out[42]: 'These are real .'

```
In [43]: eng[0]
```

Out[43]: 'These are real.'

```
In [44]: por_recoded[0]
```

Out[44]: 'Estas são autênticas .'

```
In [45]: por[0]
```

Out[45]: 'Estas são autênticas.'

```
In [46]: def encode_seq(x,tokenizer,max_length=0):
# String to integer
x = tokenizer.encode("<START>" + x + "<STOP>").ids
x += [0]*(max_length-len(x))
return x

def decode_seq(x,tokenizer):
return tokenizer.decode(x)
```

```
In [47]: temp = encode_seq(eng_recoded[0],eng_tokenizer,20)
temp
```

Out[47]: [1, 425, 140, 442, 9, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]

```
In [48]: len(temp)
```

Out[48]: 20

```
In [49]: decode_seq(temp,eng_tokenizer)
```

Out[49]: 'These are real .'


```
In [50]: temp = encode_seq(por_recoded[0],por_tokenizer,20)
temp
```

```
Out[50]: [1, 862, 229, 6063, 8, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

```
In [51]: decode_seq(temp,por_tokenizer)
```

```
Out[51]: 'Estas são autênticas .'
```

```
In [52]: max_eng = np.max([len(encode_seq(i,eng_tokenizer)) for i in eng_recoded])
max_eng
```

```
Out[52]: 12
```

```
In [53]: max_por = np.max([len(encode_seq(i,por_tokenizer)) for i in por_recoded])
max_por
```

```
Out[53]: 11
```

```
In [54]: X = np.vstack([encode_seq(x,eng_tokenizer,max_eng) for x in eng_recoded])
Y = np.vstack([encode_seq(x,por_tokenizer,max_por) for x in por_recoded])
```

```
In [55]: enc_x_train = X[:split_point]
enc_x_val = X[split_point:]
enc_x_train
```

```
Out[55]: array([[ 1, 425, 140, ..., 0, 0, 0],
 [ 1, 30, 6, ..., 0, 0, 0],
 [ 1, 30, 152, ..., 0, 0, 0],
 ...,
 [ 1, 90, 6, ..., 0, 0, 0],
 [ 1, 30, 6, ..., 0, 0, 0],
 [ 1, 1592, 21, ..., 0, 0, 0]])
```

```
In [56]: dec_x_train = Y[:,0:-1][:split_point]
dec_x_val = Y[:,0:-1][split_point:]
dec_x_train
```

```
Out[56]: array([[ 1, 862, 229, ..., 0, 0, 0],
 [ 1, 3279, 3, ..., 0, 0, 0],
 [ 1, 352, 719, ..., 0, 0, 0],
 ...,
 [ 1, 141, 416, ..., 0, 0, 0],
 [ 1, 188, 850, ..., 0, 0, 0],
 [ 1, 3776, 19, ..., 0, 0, 0]])
```

```
In [57]: dec_y_train = Y[:,1:][:split_point]
dec_y_val = Y[:,1:][split_point:]
dec_y_train
```

```
Out[57]: array([[ 862, 229, 6063, ..., 0, 0, 0],
 [3279, 3, 2, ..., 0, 0, 0],
 [ 352, 719, 8, ..., 0, 0, 0],
 ...,
 [ 141, 416, 3597, ..., 0, 0, 0],
 [ 188, 850, 8, ..., 0, 0, 0],
 [3776, 19, 2, ..., 0, 0, 0]])
```

```
In [58]: print(enc_x_train.shape)
print(dec_x_train.shape)
print(dec_y_train.shape)
```

```
(8000, 12)
(8000, 10)
(8000, 10)
```

```
In [59]: print(enc_x_val.shape)
print(dec_x_val.shape)
print(dec_y_val.shape)
```

```
(2000, 12)
(2000, 10)
(2000, 10)
```

```
In [60]: batch_size = 256
xy_train = torch.utils.data.DataLoader(list(zip(torch.Tensor(enc_x_train).long(),
                                                torch.Tensor(dec_x_train).long(),
                                                torch.Tensor(dec_y_train).long(),
                                                shuffle=True, batch_size=batch_size,
                                                num_workers=4)
xy_val = torch.utils.data.DataLoader(list(zip(torch.Tensor(enc_x_val).long(),
                                                torch.Tensor(dec_x_val).long(),
                                                torch.Tensor(dec_y_val).long()),
                                                shuffle=False, batch_size=batch_size,
                                                num_workers=4)
```

Encoder-Decoder Network

```
In [61]: class RecurrentResidual(pl.LightningModule):
    def __init__(self,
                 latent_size = 256,
                 bidirectional = False,
                 **kwargs):
        super().__init__(**kwargs)
        self.layer_norm = torch.nn.LayerNorm(latent_size)
        self.rnn_layer = torch.nn.LSTM(latent_size,
                                       latent_size // 2 if bidirectional else latent_size,
                                       bidirectional=bidirectional,
                                       batch_first=True)

    def forward(self, x):
        return x + self.rnn_layer(self.layer_norm(x))[0]
```

```
In [62]: class EncoderNetwork(pl.LightningModule):
    def __init__(self,
                 num_tokens,
                 latent_size = 256, # Use something divisible by 2
                 n_layers = 8,
                 **kwargs):
        super().__init__(**kwargs)
        self.embedding = torch.nn.Embedding(num_tokens,
                                             latent_size,
                                             padding_idx=0)

        self.dropout = torch.nn.Dropout1d(0.05) # Whole token dropped
        self.rnn_layers = torch.nn.Sequential(*[
            RecurrentResidual(latent_size, True) for _ in range(n_layers)
```

```

    ])

    def forward(self, x):
        y = x
        y = self.embedding(y)
        y = self.dropout(y)
        y = self.rnn_layers(y)[:,-1]
        return y

```

```

In [63]: class DecoderNetwork(pl.LightningModule):
    def __init__(self,
                  num_tokens,
                  latent_size = 256, # Use something divisible by 2
                  n_layers = 8,
                  **kwargs):
        super().__init__(**kwargs)
        self.embedding = torch.nn.Embedding(num_tokens,
                                             latent_size,
                                             padding_idx=0)

        # self.dropout = torch.nn.Dropout1d(0.1) # Whole token dropped
        self.linear = torch.nn.Linear(latent_size*2,
                                       latent_size)

        self.rnn_layers = torch.nn.Sequential(*[
            RecurrentResidual(latent_size, False) for _ in range(n_layers)
        ])
        self.output_layer = torch.nn.Linear(latent_size,
                                             num_tokens)

    def forward(self, x_enc, x_dec):
        y_enc = x_enc.unsqueeze(1).repeat(1, x_dec.shape[1], 1)
        y_dec = self.embedding(x_dec)
        # y_dec = self.dropout(y_dec)
        y = y_enc
        y = torch.concatenate([y_enc, y_dec], -1)
        y = self.linear(y)
        y = self.rnn_layers(y)
        y = self.output_layer(y)
        return y

```

```

In [64]: class EncDecLightningModule(pl.LightningModule):
    def __init__(self,
                  output_size,
                  **kwargs):
        super().__init__(**kwargs)
        self.mc_acc = torchmetrics.classification.Accuracy(task='multiclass',
                                                            num_classes=output_size,
                                                            ignore_index=0)

        self.cce_loss = torch.nn.CrossEntropyLoss(ignore_index=0)

    def predict(self, x):
        return torch.softmax(self(x), -1)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=0.001)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        x_enc, x_dec, y_dec = train_batch
        y_pred = self(x_enc, x_dec)

```

```

perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
acc = self.mc_acc(y_pred.permute(*perm),y_dec)
loss = self.cce_loss(y_pred.permute(*perm),y_dec)
self.log('train_acc', acc, on_step=False, on_epoch=True)
self.log('train_loss', loss, on_step=False, on_epoch=True)
return loss

# Validate used for Teacher Forcing
def validation_step(self, val_batch, batch_idx):
    x_enc, x_dec, y_dec = val_batch
    y_pred = self(x_enc, x_dec)
    perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
    acc = self.mc_acc(y_pred.permute(*perm),y_dec)
    loss = self.cce_loss(y_pred.permute(*perm),y_dec)
    self.log('val_acc', acc, on_step=False, on_epoch=True)
    self.log('val_loss', loss, on_step=False, on_epoch=True)
    return loss

# Test used for Non-Teacher Forcing
def test_step(self, test_batch, batch_idx):
    x_enc, x_dec, y_dec = test_batch
    context = self.enc_net(x_enc)
    tokens = torch.zeros_like(x_dec).long()
    tokens[:,0] = 1
    for i in range(y_dec.shape[1]-1):
        tokens[:,i+1] = self.dec_net(context, tokens).argmax(-1)[: ,i]
    y_pred = self(x_enc, tokens)
    perm = (0,-1) + tuple(range(y_pred.ndim))[1:-1]
    acc = self.mc_acc(y_pred.permute(*perm),y_dec)
    loss = self.cce_loss(y_pred.permute(*perm),y_dec)
    self.log('test_acc', acc, on_step=False, on_epoch=True)
    self.log('test_loss', loss, on_step=False, on_epoch=True)
    return loss

def predict_step(self, predict_batch, batch_idx):
    x_enc, x_dec, y_dec = test_batch
    context = self.enc_net(x_enc)
    tokens = torch.zeros_like(x_dec).long()
    tokens[:,0] = 1
    for i in range(y_dec.shape[1]-1):
        tokens[:,i+1] = self.dec_net(context, tokens).argmax(-1)[: ,i]
    y_pred = self(x_enc, tokens)
    return y_pred

```

```

In [65]: class EncDecNetwork(EncDecLightningModule):
def __init__(self,
            num_enc_tokens,
            num_dec_tokens,
            latent_size = 256, # Use something divisible by 2
            n_layers = 8,
            **kwargs):
    super().__init__(output_size=num_dec_tokens,
                    **kwargs)
    self.save_hyperparameters()
    self.enc_net = EncoderNetwork(num_enc_tokens,latent_size,n_layers)
    self.dec_net = DecoderNetwork(num_dec_tokens,latent_size,n_layers)

def forward(self, x_enc, x_dec):
    return self.dec_net(self.enc_net(x_enc), x_dec)

```

```
In [66]: enc_dec_net = EncDecNetwork(num_enc_tokens=eng_tokenizer.get_vocab_size(),
                                     num_dec_tokens=por_tokenizer.get_vocab_size())
```

Training Time

```
In [67]: logger = pl.loggers.CSVLogger("logs",
                                       name="persistence",
                                       version="encdec-0")
```

```
In [68]: trainer = pl.Trainer(logger=logger,
                              max_epochs=30,
                              enable_progress_bar=True,
                              log_every_n_steps=0,
                              enable_checkpointing=True, # Notice this here!
                              callbacks=[pl.callbacks.TQDMProgressBar(refresh_rate=50)])
```

GPU available: True (cuda), used: True
 TPU available: False, using: 0 TPU cores
 IPU available: False, using: 0 IPUs
 HPU available: False, using: 0 HPUs

```
In [70]: trainer.validate(enc_dec_net, xy_val)
```

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
 SLURM auto-requeueing enabled. Setting signal handlers.
 Validation: | 0/? [00:00<?, ?it/s]

Validate metric	DataLoader 0
val_acc	9.861325816018507e-05
val_loss	8.858915328979492

```
Out[70]: [{'val_acc': 9.861325816018507e-05, 'val_loss': 8.858915328979492}]
```

```
In [71]: trainer.test(enc_dec_net, xy_val)
```

LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
 SLURM auto-requeueing enabled. Setting signal handlers.
 Testing: | 0/? [00:00<?, ?it/s]

Test metric	DataLoader 0
test_acc	0.0
test_loss	8.799009323120117

```
Out[71]: [{'test_acc': 0.0, 'test_loss': 8.799009323120117}]
```

```
In [72]: trainer.fit(enc_dec_net, xy_train, xy_val)
```

```
/opt/conda/lib/python3.11/site-packages/lightning/pytorch/callbacks/model_checkpoint.py:639: Checkpoint directory logs/persistence/encdec-0/checkpoints exists and is not empty.
```

```
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
```

	Name	Type	Params
0	mc_acc	MulticlassAccuracy	0
1	cce_loss	CrossEntropyLoss	0
2	enc_net	EncoderNetwork	4.1 M
3	dec_net	DecoderNetwork	7.7 M

```
-----
11.8 M    Trainable params
0         Non-trainable params
11.8 M    Total params
47.086    Total estimated model params size (MB)
SLURM auto-requeueing enabled. Setting signal handlers.
```

```
Sanity Checking: | 0/? [00:00<?, ?it/s]
```

```
Training: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
Validation: | 0/? [00:00<?, ?it/s]
```

```
`Trainer.fit` stopped: `max_epochs=30` reached.
```


```
In [73]: results = pd.read_csv(logger.log_dir+"/metrics.csv")
         results
```

Out[73]:


	val_acc	val_loss	epoch	step	test_acc	test_loss	train_acc	train_loss
0	0.000099	8.858916	0	0	NaN	NaN	NaN	NaN
1	0.000099	8.858915	0	0	NaN	NaN	NaN	NaN
2	NaN	NaN	0	0	0.0	8.799009	NaN	NaN
3	0.456022	3.693235	0	31	NaN	NaN	NaN	NaN
4	NaN	NaN	0	31	NaN	NaN	0.378197	4.549609
...
58	NaN	NaN	27	895	NaN	NaN	0.893199	0.327084
59	0.670228	2.483268	28	927	NaN	NaN	NaN	NaN
60	NaN	NaN	28	927	NaN	NaN	0.891603	0.331121
61	0.671111	2.522851	29	959	NaN	NaN	NaN	NaN
62	NaN	NaN	29	959	NaN	NaN	0.894371	0.313119

63 rows × 8 columns

```
In [74]: plt.plot(results["epoch"][np.logical_not(np.isnan(results["train_loss"]))],
                results["train_loss"][np.logical_not(np.isnan(results["train_loss"]))],
                label="Training")
plt.plot(results["epoch"][np.logical_not(np.isnan(results["val_loss"]))],
                results["val_loss"][np.logical_not(np.isnan(results["val_loss"]))],
                label="Validation")
plt.legend()
plt.ylabel("CCE Loss")
plt.xlabel("Epoch")
plt.show()
```



```
In [75]: plt.plot(results["epoch"][np.logical_not(np.isnan(results["train_acc"]))],
                results["train_acc"][np.logical_not(np.isnan(results["train_acc"]))],
                label="Training")
plt.plot(results["epoch"][np.logical_not(np.isnan(results["val_acc"]))],
         results["val_acc"][np.logical_not(np.isnan(results["val_acc"]))],
         label="Validation")
plt.legend()
plt.ylabel("Accuracy")
plt.xlabel("Epoch")
plt.show()
```

Test without Teacher Forcing

```
In [76]: # Complete max_length cycles with the decoder
i = 0
enc_dec_net.to("cpu")
context = enc_dec_net.enc_net(torch.Tensor(enc_x_val[i:i+1]).long())
token = torch.zeros((1,dec_y_val.shape[1])).long()
token[0,0] = 1

for x in range(dec_y_val.shape[1]-1):
    result = enc_dec_net.dec_net(context,token).argmax(-1)
    token[0,x+1] = result[0,x]
    if result[0,x] == 2:
        break
result = token.cpu().detach().numpy()[0]
result
```

Out[76]: array([1, 5549, 514, 19, 2, 0, 0, 0, 0])

English input...

```
In [77]: decode_seq(enc_x_val[i],eng_tokenizer)
```

Out[77]: 'Can you swim ?'

Portuguese translation from network...

```
In [78]: decode_seq(result,por_tokenizer)
```

Out[78]: 'Sabes nadar ?'

Target translation from the data set...

```
In [79]: decode_seq(dec_y_val[i], por_tokenizer)
```

```
Out[79]: 'Sabe nadar ?'
```

```
In [80]: result.shape
```

```
Out[80]: (10,)
```

```
In [81]: dec_y_val.shape
```

```
Out[81]: (2000, 10)
```

Accuracy **without** teacher forcing...

Keep track of the following result for comparison later...

```
In [82]: trainer.test(enc_dec_net, xy_val)
```

```
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
SLURM auto-requeueing enabled. Setting signal handlers.
/opt/conda/lib/python3.11/site-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(
Testing: |          | 0/? [00:00<?, ?it/s]
```

Test metric	DataLoader 0
test_acc	0.4213317930698395
test_loss	6.924009323120117

```
Out[82]: [{'test_acc': 0.4213317930698395, 'test_loss': 6.924009323120117}]
```

Here's the magic word:

```
In [83]: trainer.save_checkpoint("enc_dec_net.ckpt")
```

Switch to Loading example notebook...